

In [1]: *#importing all the required packages*

```
import numpy as np
import pandas as pd
import math
import cv2
import matplotlib.pyplot as plt
import os
import seaborn as sns

from PIL import Image
from os import listdir
from os.path import isfile, join
from random import shuffle


import tensorflow as tf
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D
from keras.layers import Activation, Dropout, Flatten, Dense
from keras.utils.np_utils import to_categorical
```

In [2]: *#specifying dataset directory*

```
os.chdir('D:\\6th sem\\Gender & Age classification\\utkface_aligned_cropped\\UTKFace')
```

In [3]: `im = Image.open('25_1_2_20170104015850244.jpg.chip.jpg').resize((64,64))`  
`im`

Out[3]:



In [4]: `onlyfiles = os.listdir()`

In [5]: *#counting number of image files in onlyfiles*

```
len(onlyfiles)
```

Out[5]: 23151

In [6]: *#shuffling all the images in onlyfiles*

```
shuffle(onlyfiles)
```

In [7]: *#getting age from the name of the image*

```
age = [i.split('_')[0] for i in onlyfiles]
```

In [8]: `df = pd.DataFrame()`  
`df['onlyfiles'], df['age'] = onlyfiles, age`

In [9]: *#displaying top files present in dataframe using head*

```
df.head()
```

Out[9]:

	onlyfiles	age
0	65_1_0_20170110183018880.jpg.chip.jpg	65
1	70_1_0_20170110122250867.jpg.chip.jpg	70
2	45_1_1_20170116003204136.jpg.chip.jpg	45
3	36_0_0_20170116182401789.jpg.chip.jpg	36
4	8_1_0_20170109202804967.jpg.chip.jpg	8

In [10]: *#displaying bottom files present in dataframe using tail*

```
df.tail()
```


Out[10]:

	onlyfiles	age
23146	32_0_2_20170116182724023.jpg.chip.jpg	32
23147	65_1_3_20170109141910621.jpg.chip.jpg	65
23148	45_1_0_20170103183543298.jpg.chip.jpg	45
23149	26_1_1_20170112205853539.jpg.chip.jpg	26
23150	49_1_3_20170109132625049.jpg.chip.jpg	49

In [11]: *#displaying the no of image corresponding with the age*

```
sns.histplot(df['age'])
```

Out[11]: <Axes: xlabel='age', ylabel='Count'>



In [12]: `classes = []`

```
#defining the age groups
```

```

for i in age:
    i = int(i)
    if i <= 5:                                #babies (1-5) CLASS 0
        classes.append(0)

    if (i>5) and (i<=12):                    #Children (6-12) CLASS 1
        classes.append(1)

    if (i>12) and (i<=18):                  #Youth (12-18) CLASS 2
        classes.append(2)

    if (i>18) and (i<40):                   #ADULTS (19-40) CLASS 3
        classes.append(3)

    if (i>=40) and (i<60):                 #MIDDLE AGED (40-60) 4
        classes.append(4)

    if i>=60:                              #Very Old (>60) CLASS 5
        classes.append(5)

```

In [13]: *# CONVERT IMAGES TO VECTORS*

In [14]: `X_data = []`

*#inserting each image one by one on X\_data*

```

for file in onlyfiles:
    face = cv2.imread(file)
    face = cv2.resize(face, (64, 64) )
    X_data.append(face)

```

In [15]: *# By applying np.squeeze(), the singleton dimensions were removed (1, 3, 1) ->*  
*#The main purpose of np.squeeze() is to reduce the dimensions of an array by removing*

```

X = np.squeeze(X_data)
# X = np.array(X_data)

```

In [16]: `X.shape`

Out[16]: (23151, 64, 64, 3)

In [17]: *# normalizing data (0-1)*

```

X = X.astype('float64')
X /= 255

```

In [18]: `classes[:10]`

Out[18]: [5, 5, 4, 3, 1, 3, 0, 3, 3, 4]

In [19]: `categorical_labels = to_categorical(classes, num_classes=6)`

In [20]: `categorical_labels[:10]`

```
Out[20]: array([[0., 0., 0., 0., 0., 1.],
        [0., 0., 0., 0., 0., 1.],
        [0., 0., 0., 0., 1., 0.],
        [0., 0., 0., 1., 0., 0.],
        [0., 1., 0., 0., 0., 0.],
        [0., 0., 0., 1., 0., 0.],
        [1., 0., 0., 0., 0., 0.],
        [0., 0., 0., 1., 0., 0.],
        [0., 0., 0., 1., 0., 0.],
        [0., 0., 0., 0., 1., 0.]], dtype=float32)
```

```
In [22]: from sklearn.model_selection import train_test_split
        # Split the dataset for training, testing, and validation
        x_train, x_test, y_train, y_test = train_test_split(X, categorical_labels, test_size=0.2)
        x_train, x_valid, y_train, y_valid = train_test_split(x_train, y_train, test_size=0.1)
```

```
In [24]: from keras.preprocessing.image import ImageDataGenerator

        # Data augmentation

        datagen = ImageDataGenerator(
            rotation_range=20,
            width_shift_range=0.1,
            height_shift_range=0.1,
            shear_range=0.2,
            zoom_range=0.2,
            horizontal_flip=True,
            fill_mode='nearest'
        )
```

```
In [26]: # Build the CNN model

        model = Sequential()
        model.add(Conv2D(filters=128, kernel_size=3, padding='same', activation='relu', input_shape=(32, 32, 3)))
        model.add(MaxPooling2D(pool_size=2))
        model.add(Dropout(0.25))

        model.add(Conv2D(filters=64, kernel_size=3, padding='same', activation='relu'))
        model.add(MaxPooling2D(pool_size=2))
        model.add(Dropout(0.25))

        model.add(Conv2D(filters=32, kernel_size=3, padding='same', activation='relu'))
        model.add(MaxPooling2D(pool_size=2))
        model.add(Dropout(0.25))

        model.add(Flatten())

        model.add(Dense(512, activation='relu'))
        model.add(Dropout(0.5))
        model.add(Dense(6, activation='softmax'))

        model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 64, 64, 128)	3584
max_pooling2d_3 (MaxPooling 2D)	(None, 32, 32, 128)	0
dropout_4 (Dropout)	(None, 32, 32, 128)	0
conv2d_4 (Conv2D)	(None, 32, 32, 64)	73792
max_pooling2d_4 (MaxPooling 2D)	(None, 16, 16, 64)	0
dropout_5 (Dropout)	(None, 16, 16, 64)	0
conv2d_5 (Conv2D)	(None, 16, 16, 32)	18464
max_pooling2d_5 (MaxPooling 2D)	(None, 8, 8, 32)	0
dropout_6 (Dropout)	(None, 8, 8, 32)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_2 (Dense)	(None, 512)	1049088
dropout_7 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 6)	3078
Total params: 1,148,006		
Trainable params: 1,148,006		
Non-trainable params: 0		

In [30]: *#compiling model*

```
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

In [31]: *# Fit the model with data augmentation*

```
history = model.fit(datagen.flow(x_train, y_train, batch_size=64),
                    steps_per_epoch=len(x_train) // 64,
                    epochs=25,
                    validation_data=(x_valid, y_valid))
```

Epoch 1/30  
251/251 [=====] - 185s 725ms/step - loss: 1.3028 - accuracy: 0.5356 - val\_loss: 1.1135 - val\_accuracy: 0.6232

Epoch 2/30  
251/251 [=====] - 183s 731ms/step - loss: 1.0088 - accuracy: 0.6282 - val\_loss: 0.9303 - val\_accuracy: 0.6680

Epoch 3/30  
251/251 [=====] - 183s 731ms/step - loss: 0.9292 - accuracy: 0.6522 - val\_loss: 0.8459 - val\_accuracy: 0.6858

Epoch 4/30  
251/251 [=====] - 183s 731ms/step - loss: 0.8741 - accuracy: 0.6709 - val\_loss: 0.8013 - val\_accuracy: 0.6977

Epoch 5/30  
251/251 [=====] - 183s 731ms/step - loss: 0.8319 - accuracy: 0.6863 - val\_loss: 0.7776 - val\_accuracy: 0.7043

Epoch 6/30  
251/251 [=====] - 182s 726ms/step - loss: 0.8046 - accuracy: 0.6918 - val\_loss: 0.7284 - val\_accuracy: 0.7233

Epoch 7/30  
251/251 [=====] - 183s 730ms/step - loss: 0.7797 - accuracy: 0.6990 - val\_loss: 0.7103 - val\_accuracy: 0.7215

Epoch 8/30  
251/251 [=====] - 89s 354ms/step - loss: 0.7540 - accuracy: 0.7067 - val\_loss: 0.7249 - val\_accuracy: 0.7283

Epoch 9/30  
251/251 [=====] - 65s 258ms/step - loss: 0.7326 - accuracy: 0.7157 - val\_loss: 0.6834 - val\_accuracy: 0.7290

Epoch 10/30  
251/251 [=====] - 65s 259ms/step - loss: 0.7151 - accuracy: 0.7205 - val\_loss: 0.6787 - val\_accuracy: 0.7327

Epoch 11/30  
251/251 [=====] - 66s 264ms/step - loss: 0.7000 - accuracy: 0.7237 - val\_loss: 0.6826 - val\_accuracy: 0.7313

Epoch 12/30  
251/251 [=====] - 66s 263ms/step - loss: 0.6920 - accuracy: 0.7295 - val\_loss: 0.7028 - val\_accuracy: 0.7155

Epoch 13/30  
251/251 [=====] - 65s 259ms/step - loss: 0.6702 - accuracy: 0.7364 - val\_loss: 0.6748 - val\_accuracy: 0.7473

Epoch 14/30  
251/251 [=====] - 68s 270ms/step - loss: 0.6533 - accuracy: 0.7413 - val\_loss: 0.6659 - val\_accuracy: 0.7402

Epoch 15/30  
251/251 [=====] - 67s 269ms/step - loss: 0.6485 - accuracy: 0.7408 - val\_loss: 0.6595 - val\_accuracy: 0.7407

Epoch 16/30  
251/251 [=====] - 66s 261ms/step - loss: 0.6382 - accuracy: 0.7468 - val\_loss: 0.6585 - val\_accuracy: 0.7377

Epoch 17/30  
251/251 [=====] - 68s 271ms/step - loss: 0.6248 - accuracy: 0.7489 - val\_loss: 0.6444 - val\_accuracy: 0.7500

Epoch 18/30  
251/251 [=====] - 66s 262ms/step - loss: 0.6208 - accuracy: 0.7538 - val\_loss: 0.6529 - val\_accuracy: 0.7452

Epoch 19/30  
251/251 [=====] - 66s 262ms/step - loss: 0.6019 - accuracy: 0.7575 - val\_loss: 0.6513 - val\_accuracy: 0.7460

Epoch 20/30  
251/251 [=====] - 67s 266ms/step - loss: 0.5945 - accuracy: 0.7602 - val\_loss: 0.6524 - val\_accuracy: 0.7470

Epoch 21/30  
251/251 [=====] - 69s 274ms/step - loss: 0.5886 - accuracy: 0.7650 - val\_loss: 0.6520 - val\_accuracy: 0.7497

Epoch 22/30

```

251/251 [=====] - 68s 270ms/step - loss: 0.5799 - accurac
y: 0.7650 - val_loss: 0.6443 - val_accuracy: 0.7490
Epoch 23/30
251/251 [=====] - 73s 293ms/step - loss: 0.5685 - accurac
y: 0.7694 - val_loss: 0.6567 - val_accuracy: 0.7472
Epoch 24/30
251/251 [=====] - 76s 303ms/step - loss: 0.5624 - accurac
y: 0.7710 - val_loss: 0.6750 - val_accuracy: 0.7455
Epoch 25/30
251/251 [=====] - 75s 300ms/step - loss: 0.5623 - accurac
y: 0.7749 - val_loss: 0.6701 - val_accuracy: 0.7443
Epoch 26/30
251/251 [=====] - 70s 280ms/step - loss: 0.5465 - accurac
y: 0.7752 - val_loss: 0.6671 - val_accuracy: 0.7452
Epoch 27/30
251/251 [=====] - 70s 279ms/step - loss: 0.5325 - accurac
y: 0.7810 - val_loss: 0.6452 - val_accuracy: 0.7520
Epoch 28/30
251/251 [=====] - 68s 271ms/step - loss: 0.5280 - accurac
y: 0.7865 - val_loss: 0.6612 - val_accuracy: 0.7480
Epoch 29/30
251/251 [=====] - 67s 266ms/step - loss: 0.5191 - accurac
y: 0.7845 - val_loss: 0.6536 - val_accuracy: 0.7543
Epoch 30/30
251/251 [=====] - 66s 262ms/step - loss: 0.5163 - accurac
y: 0.7871 - val_loss: 0.6701 - val_accuracy: 0.7440
Out[31]: <keras.callbacks.History at 0x1b7a22732b0>

```

Out[31]:

```

In [32]: # Evaluating the model on test set
score = model.evaluate(x_test, y_test)

#test accuracy
print('\n', 'Test accuracy:', score[1])

```

```

48/48 [=====] - 1s 29ms/step - loss: 0.6635 - accuracy:
0.7531

```

```

Test accuracy: 0.7531106472015381

```

```

In [33]: #Labeling age groups

```

```

labels = ["BABY",      # index 0
          "CHILD",     # index 1
          "YOUTH",     # index 2
          "ADULT",     # index 3
          "MIDDLEAGE", # index 4
          "OLD",       # index 5
          ]

```

```

In [40]: y_hat = model.predict(x_test)

```

```

# Plot a random sample of 15 test images, their predicted labels and ground truth

figure = plt.figure(figsize=(15, 8))
for i, index in enumerate(np.random.choice(x_test.shape[0], size=15, replace=False)):
    ax = figure.add_subplot(3, 5, i + 1, xticks=[], yticks=[])


    # Display each image
    ax.imshow(np.squeeze(x_test[index]))
    predict_index = np.argmax(y_hat[index])
    true_index = np.argmax(y_test[index])

    # Set the title for each image

```

```
ax.set_title("{} ({}).format(labels[predict_index],
                              labels[true_index]),
            color=("green" if predict_index == true_index else "red"))
plt.show()
```

48/48 [=====] - 1s 28ms/step



```
In [41]: model.save('D:\\6th sem\\Gender & Age classification\\Age_Classification_acc75_ep30.h5')
```

```
In [ ]: #Loading model and Detecting images from model
```

```
In [1]: from tensorflow.keras.models import load_model
        from sklearn.metrics import accuracy_score
```

```
In [5]: model=load_model('D:\\6th sem\\Gender & Age classification\\Age_Classification_acc75_ep30.h5')
```

```
In [2]: # x_test[0]
```

```
In [61]: import cv2
         import numpy as np
```

```
In [62]: image=cv2.imshow('Age',x_test[0])
         cv2.waitKey(0)
```

```
Out[62]: -1
```

```
In [66]: img_path=input("Enter the image path: ")
         img=cv2.imread(img_path)

         cv2.imshow('Age',img)
         cv2.waitKey(0)
         cv2.destroyAllWindows()

         img=np.array(img)
         img=cv2.resize(img,(64,64))
         img.shape
         img=np.expand_dims(img,0)

         output=model.predict(img)
         final_out=np.argmax(output)
         print("The given image is: ",labels[final_out])
```



```
Enter the image path: D:\6th sem\Gender & Age classification\Gender_Dataset\man\face_1075.jpg
```

```
1/1 [=====] - 0s 9ms/step
```

```
The given image is: ADULT
```

```
In [ ]: face_1075
```

```
In [ ]:
```